

Advanced Go Concurrency Patterns

Sameer Ajmani
Google

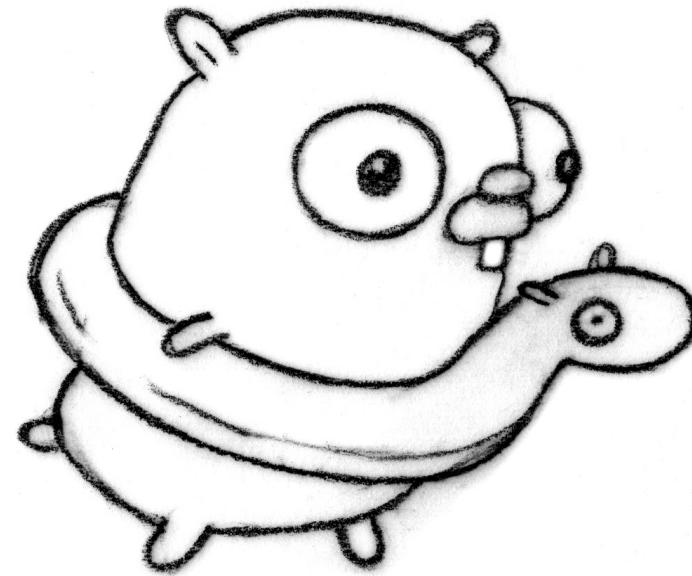
Video

This talk was presented at Google I/O in May 2013.

[Watch the talk on YouTube](https://www.youtube.com/watch?v=QDDwwePbDtw) (<https://www.youtube.com/watch?v=QDDwwePbDtw>)

2

Get ready



3

Go supports concurrency

In the language and runtime, not a library.

This changes how you structure your programs.

Goroutines and Channels

Goroutines are independently executing functions in the same address space.

```
go f()  
go g(1, 2)
```

Channels are typed values that allow goroutines to synchronize and exchange information.

```
c := make(chan int)  
go func() { c <- 3 }()  
n := <-c
```

For more on the basics, watch [Go Concurrency Patterns \(Pike, 2012\)](#) (<http://talks.golang.org/2012/concurrency.slide#1>) 5

Example: ping-pong

```
type Ball struct{ hits int }

func main() {
    table := make(chan *Ball)
    go player("ping", table)
    go player("pong", table)

    table <- new(Ball) // game on; toss the ball
    time.Sleep(1 * time.Second)
    <-table // game over; grab the ball
}

func player(name string, table chan *Ball) {
    for {
        ball := <-table
        ball.hits++
        fmt.Println(name, ball.hits)
        time.Sleep(100 * time.Millisecond)
        table <- ball
    }
}
```

Run

6

Deadlock detection

```
type Ball struct{ hits int }

func main() {
    table := make(chan *Ball)
    go player("ping", table)
    go player("pong", table)

    // table <- new(Ball) // game on; toss the ball
    time.Sleep(1 * time.Second)
    <-table // game over; grab the ball
}

func player(name string, table chan *Ball) {
    for {
        ball := <-table
        ball.hits++
        fmt.Println(name, ball.hits)
        time.Sleep(100 * time.Millisecond)
        table <- ball
    }
}
```

Run

7

Panic dumps the stacks

```
type Ball struct{ hits int }

func main() {
    table := make(chan *Ball)
    go player("ping", table)
    go player("pong", table)

    table <- new(Ball) // game on; toss the ball
    time.Sleep(1 * time.Second)
    <-table // game over; grab the ball

    panic("show me the stacks")
}

func player(name string, table chan *Ball) {
    for {
        ball := <-table
        ball.hits++
        fmt.Println(name, ball.hits)
        time.Sleep(100 * time.Millisecond)
        table <- ball
    }
}
```

Run

It's easy to go, but how to stop?

Long-lived programs need to clean up.

Let's look at how to write programs that handle communication, periodic events, and cancellation.

The core is Go's select statement: like a switch, but the decision is made based on the ability to communicate.

```
select {  
    case xc <- x:  
        // sent x on xc  
    case y := <-yc:  
        // received y from yc  
}
```

Example: feed reader

My favorite feed reader disappeared. I need a new one.

Why not write one?

Where do we start?

10

Find an RSS client

Searching godoc.org for "rss" turns up several hits, including one that provides:

```
// Fetch fetches Items for uri and returns the time when the next
// fetch should be attempted. On failure, Fetch returns an error.
func Fetch(uri string) (items []Item, next time.Time, err error)

type Item struct{
    Title, Channel, GUID string // a subset of RSS fields
}
```

But I want a stream:

```
<-chan Item
```

And I want multiple subscriptions.

11

Here's what we have

```
type Fetcher interface {  
    Fetch() ([]Item, time.Time, error)  
}  
  
func Fetch(domain string) Fetcher {...} // fetches Items from domain
```

12

Here's what we want

```
type Subscription interface {  
    Updates() <-chan Item // stream of Items  
    Close() error         // shuts down the stream  
}  
  
func Subscribe(fetcher Fetcher) Subscription {...} // converts Fetches to a stream  
  
func Merge(subs ...Subscription) Subscription {...} // merges several streams
```

Example

```
func main() {
    // Subscribe to some feeds, and create a merged update stream.
    merged := Merge(
        Subscribe(Fetch("blog.golang.org")),
        Subscribe(Fetch("googleblog.blogspot.com")),
        Subscribe(Fetch("googledevelopers.blogspot.com")))
    // Close the subscriptions after some time.
    time.AfterFunc(3*time.Second, func() {
        fmt.Println("closed:", merged.Close())
    })
    // Print the stream.
    for it := range merged.Updates() {
        fmt.Println(it.Channel, it.Title)
    }
    panic("show me the stacks")
}
```

Run

14

Subscribe

Subscribe creates a new Subscription that repeatedly fetches items until Close is called.

```
func Subscribe(fetcher Fetcher) Subscription {
    s := &sub{
        fetcher: fetcher,
        updates: make(chan Item), // for Updates
    }
    go s.loop()
    return s
}

// sub implements the Subscription interface.
type sub struct {
    fetcher Fetcher // fetches items
    updates chan Item // delivers items to the user
}

// loop fetches items using s.fetcher and sends them
// on s.updates. loop exits when s.Close is called.
func (s *sub) loop() {...}
```

Implementing Subscription

To implement the Subscription interface, define `Updates` and `Close`.

```
func (s *sub) Updates() <-chan Item {
    return s.updates
}
```

```
func (s *sub) Close() error {
    // TODO: make loop exit
    // TODO: find out about any error
    return err
}
```

What does loop do?

- periodically call Fetch
- send fetched items on the Updates channel
- exit when Close is called, reporting any error

Naive Implementation

```
for {
    if s.closed {
        close(s.updates)
        return
    }
    items, next, err := s.fetcher.Fetch()
    if err != nil {
        s.err = err
        time.Sleep(10 * time.Second)
        continue
    }
    for _, item := range items {
        s.updates <- item
    }
    if now := time.Now(); next.After(now) {
        time.Sleep(next.Sub(now))
    }
}
```

Run

```
func (s *naiveSub) Close() error {
    s.closed = true
    return s.err
}
```

18

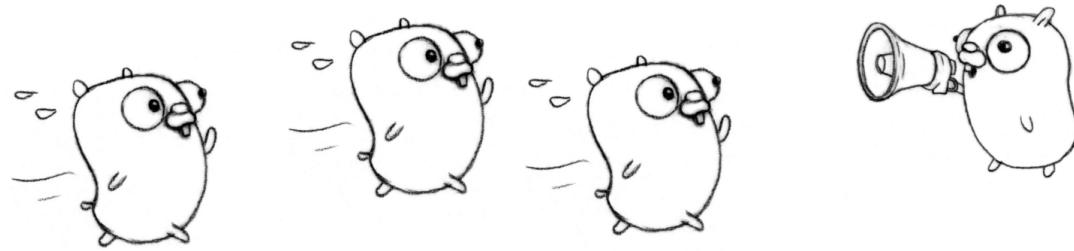
Bug 1: unsynchronized access to s.closed/s.err

```
for {
    if s.closed {
        close(s.updates)
        return
    }
    items, next, err := s.fetcher.Fetch()
    if err != nil {
        s.err = err
        time.Sleep(10 * time.Second)
        continue
    }
    for _, item := range items {
        s.updates <- item
    }
    if now := time.Now(); next.After(now) {
        time.Sleep(next.Sub(now))
    }
}
```

```
func (s *naiveSub) Close() error {
    s.closed = true
    return s.err
}
```

Race Detector

```
go run -race naivemain.go
```



```
for {
    if s.closed {
        close(s.updates)
        return
    }
    items, next, err := s.fetcher.Fetch()
    if err != nil {
        s.err = err
    }
}
```

Run

```
func (s *naiveSub) Close() error {
    s.closed = true
    return s.err
}
```

20

Bug 2: time.Sleep may keep loop running

```
for {
    if s.closed {
        close(s.updates)
        return
    }
    items, next, err := s.fetcher.Fetch()
    if err != nil {
        s.err = err
        time.Sleep(10 * time.Second)
        continue
    }
    for _, item := range items {
        s.updates <- item
    }
    if now := time.Now(); next.After(now) {
        time.Sleep(next.Sub(now))
    }
}
```

Bug 3: loop may block forever on s.updates

```
for {
    if s.closed {
        close(s.updates)
        return
    }
    items, next, err := s.fetcher.Fetch()
    if err != nil {
        s.err = err
        time.Sleep(10 * time.Second)
        continue
    }
    for _, item := range items {
        s.updates <- item
    }
    if now := time.Now(); next.After(now) {
        time.Sleep(next.Sub(now))
    }
}
```

Solution

Change the body of loop to a select with three cases:

- Close was called
- it's time to call Fetch
- send an item on s.updates

Structure: for-select loop

loop runs in its own goroutine.

select lets loop avoid blocking indefinitely in any one state.

```
func (s *sub) loop() {  
    ... declare mutable state ...  
    for {  
        ... set up channels for cases ...  
        select {  
            case <-c1:  
                ... read/write state ...  
            case c2 <- x:  
                ... read/write state ...  
            case y := <-c3:  
                ... read/write state ...  
        }  
    }  
}
```

The cases interact via local state in loop.

24

Case 1: Close

Close communicates with loop via s.closing.

```
type sub struct {  
    closing chan chan error  
}
```

The service (loop) listens for requests on its channel (s.closing).

The client (Close) sends a request on s.closing: *exit and reply with the error*

In this case, the only thing in the request is the *reply channel*.

25

Case 1: Close

Close asks loop to exit and waits for a response.

```
func (s *sub) Close() error {
    errc := make(chan error)
    s.closing <- errc
    return <-errc
}
```

loop handles Close by replying with the Fetch error and exiting.

```
var err error // set when Fetch fails
for {
    select {
    case errc := <-s.closing:
        errc <- err
        close(s.updates) // tells receiver we're done
        return
    }
}
```

Case 2: Fetch

Schedule the next Fetch after some delay.

```
var pending []Item // appended by fetch; consumed by send
var next time.Time // initially January 1, year 0
var err error
for {
    var fetchDelay time.Duration // initially 0 (no delay)
    if now := time.Now(); next.After(now) {
        fetchDelay = next.Sub(now)
    }
    startFetch := time.After(fetchDelay)

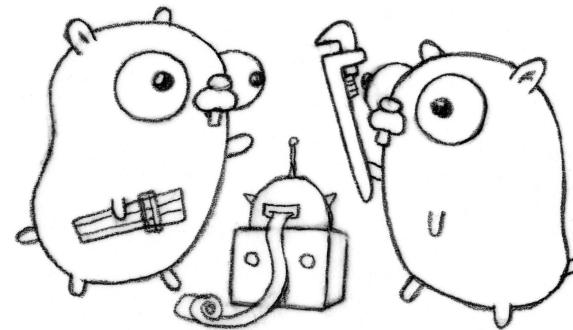
    select {
    case <-startFetch:
        var fetched []Item
        fetched, next, err = s.fetcher.Fetch()
        if err != nil {
            next = time.Now().Add(10 * time.Second)
            break
        }
        pending = append(pending, fetched...)
    }
}
```

Case 3: Send

Send the fetched items, one at a time.

```
var pending []Item // appended by fetch; consumed by send
for {
    select {
    case s.updates <- pending[0]:
        pending = pending[1:]
    }
}
```

Whoops. This crashes.



28

Select and nil channels

Sends and receives on nil channels block.

Select never selects a blocking case.

```
func main() {
    a, b := make(chan string), make(chan string)
    go func() { a <- "a" }()
    go func() { b <- "b" }()
    if rand.Intn(2) == 0 {
        a = nil
        fmt.Println("nil a")
    } else {
        b = nil
        fmt.Println("nil b")
    }
    select {
    case s := <-a:
        fmt.Println("got", s)
    case s := <-b:
        fmt.Println("got", s)
    }
}
```

Run

29

Case 3: Send (fixed)

Enable send only when pending is non-empty.

```
var pending []Item // appended by fetch; consumed by send
for {
    var first Item
    var updates chan Item
    if len(pending) > 0 {
        first = pending[0]
        updates = s.updates // enable send case
    }

    select {
    case updates <- first:
        pending = pending[1:]
    }
}
```

30

Select

Put the three cases together:

```
select {
    case errc := <-s.closing:
        errc <- err
        close(s.updates)
        return
    case <-startFetch:
        var fetched []Item
        fetched, next, err = s.fetcher.Fetch()
        if err != nil {
            next = time.Now().Add(10 * time.Second)
            break
        }
        pending = append(pending, fetched...)
    case updates <- first:
        pending = pending[1:]
}
```

The cases interact via `err`, `next`, and `pending`.

No locks, no condition variables, no callbacks.

31

Bugs fixed

- Bug 1: unsynchronized access to s.closed and s.err
- Bug 2: time.Sleep may keep loop running
- Bug 3: loop may block forever sending on s.updates

```
select {
case errc := <-s.closing:
    errc <- err
    close(s.updates)
    return
case <-startFetch:
    var fetched []Item
    fetched, next, err = s.fetcher.Fetch()
    if err != nil {
        next = time.Now().Add(10 * time.Second)
        break
    }
    pending = append(pending, fetched...)
case updates <- first:
    pending = pending[1:]
}
```

We can improve loop further

33

Issue: Fetch may return duplicates

```
var pending []Item
var next time.Time
var err error
```

```
case <-startFetch:
    var fetched []Item
    fetched, next, err = s.fetcher.Fetch()
    if err != nil {
        next = time.Now().Add(10 * time.Second)
        break
    }
    pending = append(pending, fetched...)
```

Fix: Filter items before adding to pending

```
var pending []Item
var next time.Time
var err error
var seen = make(map[string]bool) // set of item.GUIDs
```

```
case <-startFetch:
    var fetched []Item
    fetched, next, err = s.fetcher.Fetch()
    if err != nil {
        next = time.Now().Add(10 * time.Second)
        break
    }
    for _, item := range fetched {
        if !seen[item.GUID] {
            pending = append(pending, item)
            seen[item.GUID] = true
        }
    }
}
```

Issue: Pending queue grows without bound

```
case <-startFetch:  
    var fetched []Item  
    fetched, next, err = s.fetcher.Fetch()  
    if err != nil {  
        next = time.Now().Add(10 * time.Second)  
        break  
    }  
    for _, item := range fetched {  
        if !seen[item.GUID] {  
            pending = append(pending, item)  
            seen[item.GUID] = true  
        }  
    }  
}
```

Fix: Disable fetch case when too much pending

```
const maxPending = 10

var fetchDelay time.Duration
if now := time.Now(); next.After(now) {
    fetchDelay = next.Sub(now)
}
var startFetch <-chan time.Time
if len(pending) < maxPending {
    startFetch = time.After(fetchDelay) // enable fetch case
}
```

Could instead drop older items from the head of pending.

37

Issue: Loop blocks on Fetch

```
case <-startFetch:  
    var fetched []Item  
    fetched, next, err = s.fetcher.Fetch()  
    if err != nil {  
        next = time.Now().Add(10 * time.Second)  
        break  
    }  
    for _, item := range fetched {  
        if !seen[item.GUID] {  
            pending = append(pending, item)  
            seen[item.GUID] = true  
        }  
    }  
}
```

Fix: Run Fetch asynchronously

Add a new select case for fetchDone.

```
type fetchResult struct{ fetched []Item; next time.Time; err error }
```

```
var fetchDone chan fetchResult // if non-nil, Fetch is running
```

```
var startFetch <-chan time.Time
if fetchDone == nil && len(pending) < maxPending {
    startFetch = time.After(fetchDelay) // enable fetch case
}
```

```
select {
case <-startFetch:
    fetchDone = make(chan fetchResult, 1)
    go func() {
        fetched, next, err := s.fetcher.Fetch()
        fetchDone <- fetchResult{fetched, next, err}
    }()
case result := <-fetchDone:
    fetchDone = nil
    // Use result.fetched, result.next, result.err
```

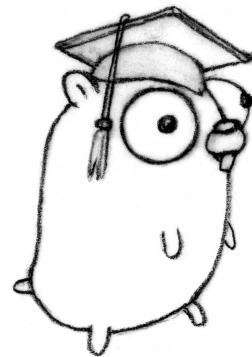
Implemented Subscribe

Responsive. Cleans up. Easy to read and change.

Three techniques:

- for-select loop
- service channel, reply channels (chan chan error)
- nil channels in select cases

More details online, including Merge.



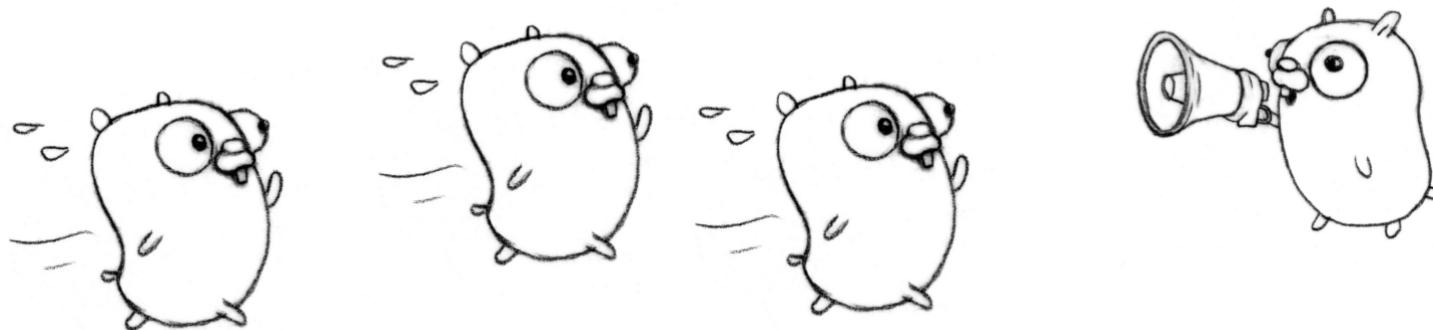
40

Conclusion

Concurrent programming can be tricky.

Go makes it easier:

- channels convey data, timer events, cancellation signals
- goroutines serialize access to local mutable state
- stack traces & deadlock detector
- race detector



41

Links

Go Concurrency Patterns (2012)

talks.golang.org/2012/concurrency.slide (<http://talks.golang.org/2012/concurrency.slide>)

Concurrency is not parallelism

golang.org/s/concurrency-is-not-parallelism (<http://golang.org/s/concurrency-is-not-parallelism>)

Share memory by communicating

golang.org/doc/codewalk/sharemem (<http://golang.org/doc/codewalk/sharemem>)

Go Tour (learn Go in your browser)

tour.golang.org (<http://tour.golang.org>)

42

Thank you

Sameer Ajmani

Google

<http://profiles.google.com/ajmani> (<http://profiles.google.com/ajmani>)

@Sajma (<http://twitter.com/Sajma>)

<http://golang.org> (<http://golang.org>)

